



WHAT'S IN A NAME? NAMING BIG SCIENCE DATA IN NAMED DATA NETWORKING

Susmit Shannigrahi, Chengyu Fan, Craig Partridge

sshannigrahi@tntech.edu




Names and the Network



<https://xkcd.com/1226/>
<https://1226.xkcd.com>


151.101.192.67




Individual apps

Every node

Individual links





Names in Named Data Networking



Individual apps

Every node

Individual links



/xkcd/1226/
/1226/xkcd
/1227/xkcd
/1228/xkcd

FIB →
/xkcd/
/1226
/1227
/1228

Motivation

- Naming is *the* central construct for most NDN operations
- In NDN, naming is intertwined with network and application performance
- Lack of systematic study on NDN naming

Goal

- What names make sense in the context of scientific applications?
 - Can we create them from existing names?
- What are the trade-offs of such names?
- A number of scientific communities can benefit from NDN
 - Can we come up with a list of naming recommendations for future applications?

Outline

- Motivation
- **Existing Names and Translating them in NDN**
- How does naming affect NDN operations?
- Naming recommendations

Existing Names in Scientific Communities

- Follows community driven hierarchical naming schemes
 - Community driven
 - Strict enforcement is often lacking
- Follows directory structure for data management
 - Files need to be identified as part of a larger collection
- Names have domain specific information
 - needed for applications and users

A Climate File Name

Original Name



http://esgf-data1.diasjp.net/thredds/fileServer/esg_dataroot/cmip5/output1/MIROC/MIROC4h/historical/6hr/atmos/6hrPlev/r1i1p1/v20110926/psl/psl_6hrPlev_MIROC4h_historical_r1i1p1_1984010100-1984123118.nc



/CMIP5/output/MIROC/MIROC4h
/historical/6hr/atmos/psl/r1i1p1/1950010100-1950123118/



NDN Name

A Root File Name (Physics)

Original Name



root://eospublic.cern.ch//eos/opendata/cms/MonteCarlo
Upgrade/RunIIFall18wmLHEGS/TTToHadronic TuneCP5_
13TeV-powheg-pythia8/GEN-SIM/IdealConditions_102X_
upgrade2018_design_v9_ext1-v2/260000/
4FF18384-7B96-9E4A-B611-AA5D176A4156.root



/cms/MonteCarloUpgrade/RunIIFall18wm
LHEGS/TTToHadronic_TuneCP5_13TeV-powheg-pythia8
/GEN-SIM/IdealConditions_102X_upgrade2018_design_v9_
ext1-v2/260000/<UUID.root>

NDN Name

A Genome Name



Original Name



<https://trace.ncbi.nlm.nih.gov/Traces/sra/?run=SRR5139424>



/BIOLOGY/SRA/9605/9606/NaN/RNA-Seq/ILLUMINA/
TRANSCRIPTOMIC/PAIRED/Kidney/PRJNA359795/
SRP095950/SRX2458154/SRR5139424/SRR5139424_1



NDN Name

The takeaways

- The concept of hierarchy exists in scientific communities
 - Some communities have already named their data
 - Others can convert their names into these *community defined* hierarchies
- Good news for NDN!

Outline

- Motivation
- Existing Names and Translating them in NDN
- **How does naming affect NDN operations?**
- Naming recommendations

NDN Names – Unresolved Questions

- Possible to create hierarchical and semantically meaningful names
- How much information should be included?
- In which order?
- How should the network handle them?
- What names should the applications use?
- Which prefixes should be announced into the network?

Expressive vs Shorter names

- A design continuum rather than exact lengths
- Expressive names carry more information
 - Encodes more domain specific information
 - Provide better context for both the network and applications
 - e.g., probabilistic caching → cache any content that has “cmip5” in the name
 - Forwarding → More granular namespaces
- Shorter names are better for reducing in-network states
 - Less useful for providing context
 - Less overhead (generally)
 - Smaller in-network state

Expressive vs Shorter names

/genus/species/subspecies

	Expressive Names	Shorter Names
Caching	(=) Popularity Based (+) Better context for probabilistic/Predictive/Context aware caching	(=) Popularity Based
Forwarding	(+) Fine grained forwarding (-) Can be slower (-) Overhead is higher	(+) Less State (+) Potentially faster (-) Coarse granularity
Routing	(+) Better route choices (-) More announcements	(+) Fewer announcements (-) Producer has less control over forwarding

Expressive vs Shorter names – Data Structures

/genus/species/subspecies/.../

	Expressive Names	Shorter Names
Interest/Data Packets	(-) Larger, less space for payloads	(=) Popularity Based
FIB/RIB	(-) Larger	(+) Smaller
Measurement Tables	(-) Larger	(+) Smaller

Expressive vs Shorter names – Application Layer Functionality

/genus/species1/subspecies
/genus/species2/subspecies

	Expressive Names	Shorter Names
Publication	(+) Needed for partial replication (-) Can use forwarding hints	(-) Difficult to support partial replication
Data Discovery	(+) More context is better for discovery and search	(-) Less context
Privacy	(-) Exposes more information	(+) Exposes less information

Decision Time

- Expressive names are better suited for both network operations and applications
- The trade-offs are: smaller payload and larger in-network state
- Which one should the scientific communities prefer?
 - Depends on the application and the network operator
 - We are only pointing out the trade-offs

Minimum Usable Data Units (MDU)

- In science communities, a set of files are often used together
 - Single files are useless
- Makes sense to treat them as a single entity
 - the Interest and Data packets should be handled consistently (e.g., same upstream and downstream, same strategy)
 - Cached together and subsequent requests sent past it
- **We are not proposing prefetching**
 - A hint to the network layer, not a requirement
 - Can be implemented as a strategy

```
/Trichosporon/asahii/1.1/ht2
/Trichosporon/asahii/1.2/ht2
/Trichosporon/asahii/1.3/ht2
/Trichosporon/asahii/1.4/ht2
/Trichosporon/asahii/1.5/ht2
/Trichosporon/asahii/1.6/ht2
/Trichosporon/asahii/1.7/ht2
/Trichosporon/asahii/1.8/ht2
/Trichosporon/asahii/1/fa
/Trichosporon/asahii/1/gff3
/Trichosporon/asahii/1/gtf
/Trichosporon/asahii/1/meta.json
/Trichosporon/asahii/1/Splice_sites
```

Outline

- Motivation
- Existing Names and Translating them in NDN
- How does naming affect NDN operations?
- Naming recommendations

Naming Recommendations for Applications

- Constructed using well-defined, hierarchical, human-readable, and semantically meaningful components
 - Provides more context to the users, applications, and network operators
- Longer names are more useful
 - Trade-off is space, can use linked objects for both shorter and longer names
- Don't use location information as part of the name
- Allow the names to be identified as part of the MDU
- For larger namespaces, use forwarding hints
- Put any sensitive information towards the end and encrypt it

Naming Recommendations for Science Network Operators

- Shorter names help with forwarding speed (generally), expressive names help with everything else (caching, forwarding)
- Data names should be location agnostic
 - */cmip5/csu* is easier to handle in NDN than */csu/cmip5*
 - Unless you need to know the location
 - e.g., Load-balancing between multiple sources
- Cache and forward MDUs together for better efficiency

Conclusions

- Naming is the central construct of NDN
- Efficient naming can align named data with application requirements and network characteristics
- For science data, we define the concept of MDU
 - handle a set of files together
 - benefits both the application and the network
- We provide a list of recommendations
 - Up to the community and the operators to implement them